

CS4055

High Performance Data Networking

Project Report

Comparative Study of P4 and OpenFlow monitoring application

Group 2

CS4055 High Performance Data Networking Project Report

**Comparative Study of P4 and OpenFlow
monitoring application**

March 22, 2021

Contents

1	Abstract	3
2	OpenFlow Monitoring Implementation & Testing	4
2.1	OpenFlow Monitoring	4
2.2	Bandwidth Monitoring.....	5
2.2.1	Implementation.....	5
2.2.2	Testing.....	6
2.3	Delay Monitoring.....	10
2.3.1	Implementation.....	10
2.3.2	Testing.....	11
2.4	Loss Monitoring.....	15
2.4.1	Implementation.....	15
2.4.2	Testing.....	16
2.5	Comparison.....	19
2.5.1	Advantages of Active approach	19
2.5.2	Disadvantages of Active approach.....	19
3	P4 Monitoring Implementation & Testing	20
3.1	P4 Monitoring	20
3.2	Diagnostic Node & Cloning.....	21
3.3	Delay Monitoring.....	23
3.3.1	Implementation.....	23
3.3.2	Testing.....	24
3.4	Bandwidth Monitoring.....	27
3.4.1	Implementation.....	27
3.4.2	Testing.....	29
3.5	Loss Monitoring.....	32
3.5.1	Implementation.....	32
3.5.2	Testing.....	32
4	Results & Comparison	37
4.1	Openflow.....	37
4.1.1	Bandwidth monitoring.....	37

4.1.2	Delay Monitoring	38
4.1.3	Loss Monitoring	38
4.2	P4	38
4.2.1	Bandwidth monitoring	38
4.2.2	Delay Monitoring	39
4.2.3	Loss Monitoring	39
4.3	OpenFlow and P4 Comparison	40
4.3.1	Openflow monitoring approach	40
4.3.2	P4 monitoring approach	40
4.3.3	Comparison based on Overhead	40
4.3.4	Comparison based on Network Statistics	41
5	Conclusion	42
	Bibliography	43

Achievements

OpenFlow

- Extended **Monitoring_Routing_Bandwidth_Delay.py**¹ with an accurate delay monitoring function and a packet loss monitoring function.
- Bandwidth Monitoring: Tested and analyzed the bandwidth monitoring function with different network configurations including different bandwidth settings, different topology and multiple simultaneous flows.
- Delay Monitoring: Calculated the delay by adding additional flow request message to remove the round trip time between the OpenFlow controller and switches. Tested and analyzed the delay monitoring function with different network configurations including different delay settings, different topology.
- Loss Monitoring: Calculated the packet loss by the returned statistics from flow requests messages and probe message. Tested and analyzed the loss monitoring function with different network configurations including different loss settings, different topology and different protocols.

P4

- Completed the basic P4 network forwarding functionality
- Delay Monitoring: show which switch(es) in the network is slowest and is the bottle neck in a network path
- Bandwidth Monitoring: monitor amount of packets on a path and reduce overloaded traffics to maintain bandwidth balance in a network
- Loss Monitoring: show at which switches most packets are dropped in the network path
- P4 monitoring fully with SimpleSwitchGrpc and gRPC controller with P4Runtime support instead of SimpleSwitch
- Using experimental Thrift server feature to enable BMv2 debugger and runtime_CLI

Contribution

- OpenFlow Implementation (Hang)
- OpenFlow Testing (Hang & Shubhra)
- P4 Implementation proposal (Shubhra)

¹This is the code example provided by the 2020-2021 Reader Solution for bandwidth monitoring and delay monitoring of the OpenFlow switch

- P4 Implementation (Tat)
- P4 Testing (Tat & Shubhra)
- Results Analysis (Tat, Hang, Shubhra)
- Comparison of P4 and OpenFlow (Shubhra)

1

Abstract

A network monitoring application collects network statistics such as propagation delay between source node and destination node, number of packets lost while traversing the route and bandwidth utilisation of the network. These applications help network designers in optimising network capacity, provide security analysis as well as give insights about the traffic which passes over the network[3]. A network monitoring application has been developed for both OpenFlow as well as P4 which effectively collects network data. This data will be used for comparing OpenFlow and P4 on the basis of the above stated network characteristics. This comparison will benefit the operators, who might be interested in implementing OpenFlow and P4.

2

OpenFlow Monitoring Implementation & Testing

2.1. OpenFlow Monitoring

In earlier networks, there was no separation of control plane and data plane which made those networks inefficient as adding new rules and updates was a challenge. With the emergence of Software Defined Networks(SDN), the data plane and control plane can be maintained virtually and function independent of each other[2]. OpenFlow is a communication protocol which implements this logical separation of data plane and control plane. In OpenFlow, switches are dumb forwarding devices and controller is responsible for installing all the rules. Using OpenFlow switch protocols these rules can be added, deleted and updated.

The OpenFlow Monitoring application is implemented on the basis of the monitoring application provided by the project reader of **2020-2021 CS4055 High Performance Data Networking**. The code of our application extends **Monitoring_Routing_Bandwidth_Delay.py**¹ with an accurate delay monitoring function and a packet loss function.

The OpenFlow Monitoring application is developed with the instant virtual network **Mininet** and the OpenFlow controller framework **Ryu**. The application is able to monitor the bandwidth, delay, and packet loss for data flows in a network. The system architecture of our application is shown in Figure 2.13. The control plane is able to route the data flow through network devices and collect QoS statistics. The variables of the monitoring application is declared in Table 2.1

¹This is the code example provided by the 2020-2021 Reader Solution for bandwidth monitoring and delay monitoring of the OpenFlow switch

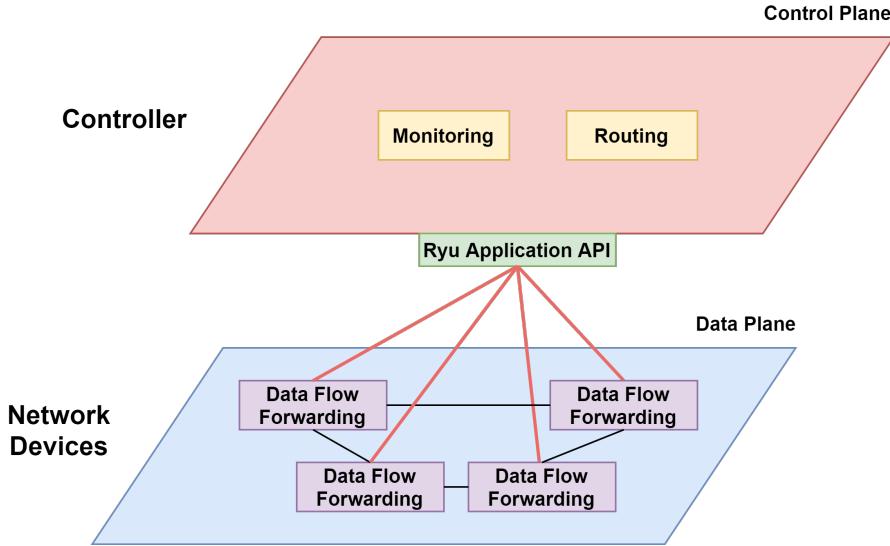


Figure 2.1: System Architecture

2.2. Bandwidth Monitoring

2.2.1. Implementation

The bandwidth of a data flow can be obtained by sending the **OFPFlowStatsRequest** message to the destination switch in the data path of the monitored data flow periodically (every **t_int**). It is presented in Figure 2.2.

variable	description
destinationSwitch	the destination switch in the data flow
sourceSwitch	the source switch in the data flow
src_mac	the mac address of the source switch in the data flow
dst_mac	the mac address of the destination switch in the data flow
sourcePort	the port from the first switch which the probe packet will be sent out to
delay	delay of the probe packet
throughput	the throughput of the link between the source switch and the destination switch
bytes_last	the total bytes of the data flow at the previous timestamp
rtt_s_c_src	the round trip time between the source switch and the controller
rtt_s_c_dst	the round trip time between the destination switch and the controller
packet_src_last	the total packets sent by the source switch at the previous timestamp
packet_dst_last	the total packets received by the destination switch at the previous timestamp
packet_src	the total packets sent by the source switch during the polling cycle
packet_dst	the total packets received by the destination switch during the polling cycle
packet_loss	the packet loss during the polling cycle
src_timestamp	the current timestamp for the request message of the source switch
dst_timestamp	the current timestamp for the request message of the destination switch

Table 2.1: variables declaration for OpenFlow Monitoring Application

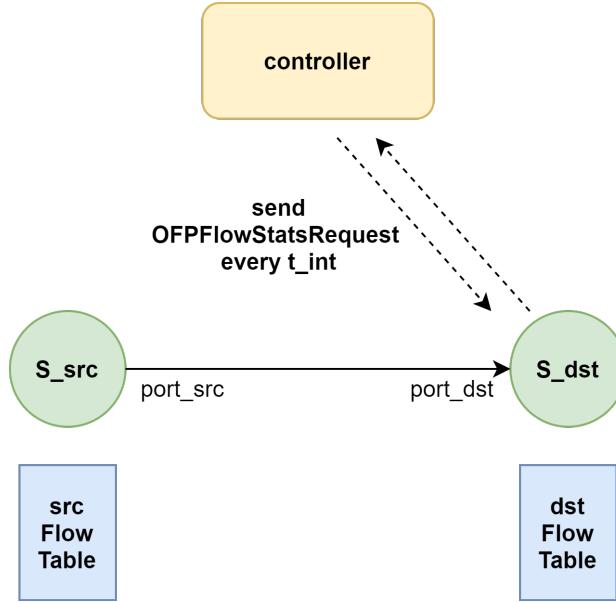


Figure 2.2: Bandwidth Monitoring

2.2.2. Testing

To verify the functionality of bandwidth monitoring, we test the application with different bandwidth settings with different topology as well as varying client host configurations. Linear topology and ring topology are used in tests. Examples of the two topology are shown in Figure 2.3 and Figure 2.4. The example command for those tests is shown below:

```
sudo mn --mac --arp --topo=linear,5 --controller=remote \
--switch=ovs, protocol=OpenFlow13 \
--links=tc,bw=20
```

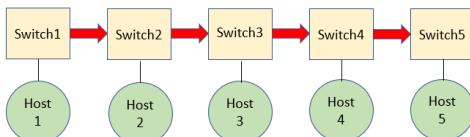


Figure 2.3: Linear Topology

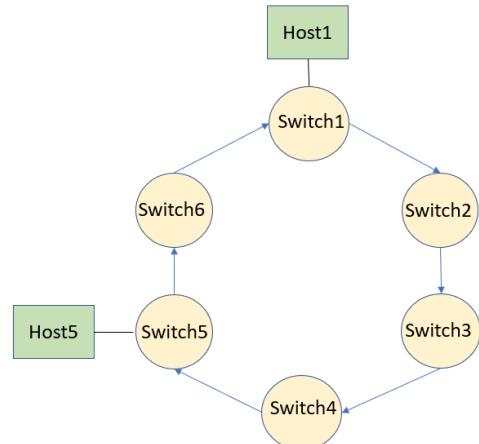


Figure 2.4: Ring Topology

Table 2.2 gives an overview of all the tests for bandwidth monitoring. In Figure 2.7 and Figure 2.11, the client of server h1 is h2, the client of server h4 is h5.

2.3. Delay Monitoring

2.3.1. Implementation

The Delay Monitoring applies the **active method** of QoS monitoring, which means it injects additional probe packets to monitor the network behaviour [7]. In order to calculate the delay of a data flow, we inject a delay probe packet to the network, the data path of the delay probe packet is the same as the data flow which is monitored. It is demonstrated in Figure 2.13. t_{probe} is the duration of the delay probe packet travelling in the data path. t_1 and t_3 are the duration of the delay probe packet traveling between controller and the source switch or the destination switch. Thus, the delay of the data flow between the source switch and the destination switch can be calculated as:

$$t_{delay} = t_{probe} - t_1 - t_3 \quad (2.1)$$

By processing the **PacketIn** event message, t_{probe} is able to obtained by the current time subtracting the time of the controller sent the delay probe packet.

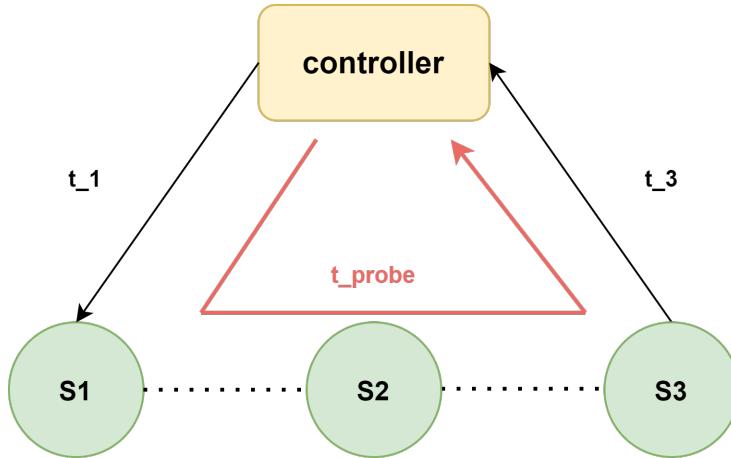


Figure 2.13: Delay Monitoring

t_1 or t_3 can be regarded as half of the round trip time between the controller and the source switch or the destination switch. The round trip time can be obtained by the time different between the **OFPFlowStatsRequest** and the **OFPFlowStatsReply**. It is demonstrated in Figure 2.14. In order to record the time t_a of sending the request message to the aimed switch, a timestamp is sent along with **OFPFlowStatsRequest**. Thus, the delay can be calculated as:

$$t_{delay} = t_{probe} - 0.5 \times (t_{rtt_src} + t_{rtt_dst}) \quad (2.2)$$

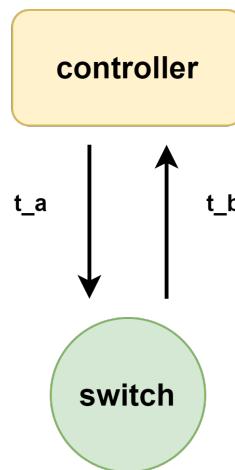


Figure 2.14: Round Trip Time

2.3.2. Testing

To verify the functionality of delay monitoring, we test the application with different delay settings and different topology configurations (**linear topology**, **complete graph topology**). Since **t_1** and **t_3** are too small to notice, we inject a delay of 10ms in the loopback interface in the network.

```
sudo tc qdisc add dev lo root netem delay 10ms
```

Table 2.3 gives an overview of all the tests for delay monitoring.

Figure	Topology	Delay (ms)	Server	Client
2.15	Linear	15	h1	h2
2.16	Linear	15	h1	h5
2.17	Linear	15	h1, h4	h2, h5
2.18	Complete Graph	15	h1	h5
2.19	Linear	30	h1	h2
2.20	Linear	30	h1	h5
2.21	Linear	30	h1, h4	h2, h5
2.22	Complete Graph	30	h1	h2

Table 2.3: Delay Monitoring Tests

the code section presented in Figure 2.23 of OpenFlow.py below, it requests statistics of the source switch and the destination switch from the OpenFlow controller and sends the latency probe packets by the OpenFlow controller to capture **t_probe**. The statistics for monitored paths are stored in the corresponding flow class instances presented in Figure 2.24. However, sometimes the update of variables in a flow class instances are not simultaneous due to the time processing the request message. Thus, there are some stacked intermediate variables in delay calculation.

```

140     def _monitor(self):
141         """
142             Main method for the monitoring actions.
143         """
144         while True:
145             self.logger.info('Monitoring stats: num_flows:' + str(len(self.monitored_paths)))
146             for m in self.monitored_paths:
147                 print m
148                 # send request statistics command to the source switch of the monitored flow
149                 self.request_stats(m.sourceSwitch, m.src_mac, m.dst_mac, time.time())
150                 # send request statistics command to the destination switch of the monitored flow
151                 self.request_stats(m.destinationSwitch, m.src_mac, m.dst_mac, time.time())
152                 self.send_latency_probe_packet(m.sourceSwitch, m.sourcePort, m.src_mac, m.dst_mac)
153             hub.sleep(self.monitor_period)

```

Figure 2.23: Method **_monitor**

```

24 class Flow():
25     def __init__(self, src, dst, sid, did, port):
26         #values needed to identify the flow (src mac, dst mac, first and last switch)
27         self.destinationswitch = did
28         self.sourceSwitch = sid
29         self.src_mac = src
30         self.dst_mac = dst
31         #port from the first switch specifying the port to which the probe packet will be sent out to
32         self.sourcePort = port
33         self.delay = 0 #delay of the probe packet
34         self.throughput = 0
35         self.bytes_last = 0
36         ##### self.rtt_s_c_src = 0 #the round trip time between the source switch and the controller
37         ##### self.rtt_s_c_dst = 0 #the round trip time between the destination switch and the controller
38
39         self.packet_src_last = 0
40         self.packet_dst_last = 0
41         ##### self.packet_src = 0 #the total packet sent by the source switch
42         ##### self.packet_dst = 0 #the total packet received by the dst switch
43
44         self.packet_loss = 0 #packet loss
45
46         self.src_timestamp = 0 #record the current timestamp of the request message of source switch
47         self.dst_timestamp = 0 #record the current timestamp of the request message of dst switch
48
49         self.monitor_period = 0 #record the time interval for monitoring
50

```

Figure 2.24: Class **flow**

2.4. Loss Monitoring

2.4.1. Implementation

The loss monitoring is similar to the implementation of bandwidth monitoring. The loss of the data flow in this project is calculated as:

$$\text{packet_loss} = (\text{packet_sent} - \text{packet_received}) / t_{\text{int}} \quad (2.3)$$

The counts of packets sent by the source switch and the counts of packets sent by the destination switch can be obtained by the **OFPFlowStatsRequest**. It is demonstrated in the Figure 2.25.

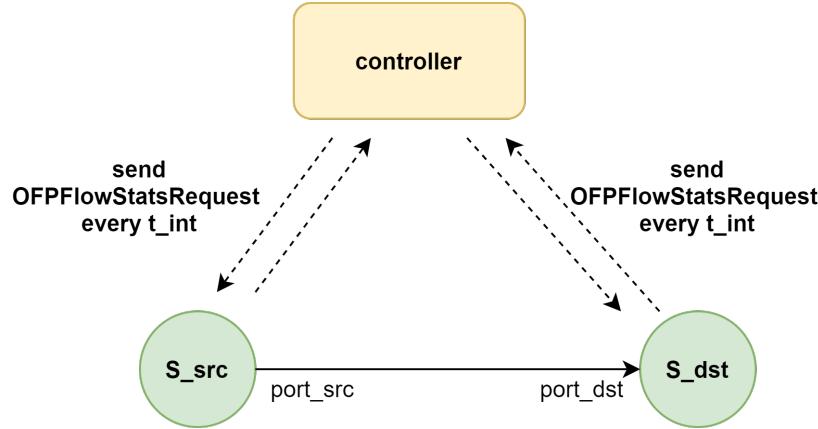


Figure 2.25: Loss Monitoring

2.4.2. Testing

To verify the functionality of loss monitoring, we test the application with different loss settings and different protocols (TCP, UDP). An example command is shown below:

```
sudo mn --mac --arp --topo=linear,5 --controller=remote \
--switch=ovs,protocol=OpenFlow13 \
--links=tc,bw=20,delay=15ms,loss=5
```

Table 2.4 gives an overview of all the tests for loss monitoring.

Figure	Topology	Loss Rate (%)	Server	Client	Protocol
2.26	Linear	0	h1	h2	tcp
2.27	Linear	0	h1	h2	udp
2.28	Linear	5	h1	h2	tcp
2.29	Linear	5	h1	h2	udp
2.30	Complete Graph	0	h1	h5	udp
2.31	Complete Graph	5	h1	h5	udp

Table 2.4: Loss Monitoring Tests

```

1 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:02) - (0 kbps 0.0 ms) (0 ms 0 ns 0 ms) (0 0 0 pkt/s 0 %)
2 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (0 kbps 0.0 ms) (0 ms 0 ns 0 ms) (0 0 0 pkt/s 0 %)
3 s1-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (1 kbps 0.055372238 ms) (17.9491043691 ms 2.94613838196 ms 2.84899578857 ms) (6 6 0 pkt/s 0 %)
4 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (2 kbps 0.03060426483 ms) (25.0788582426 ms 2.36797332764 ms 1.72805786133 ms) (7 7 0 pkt/s 0 %)
5 s2-->s2 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (94 kbps 14.7584676743 ms) (17.1530246735 ms 2.56609916659 ms 2.22301487154 ms) (362 359 3 pkt/s 0 %)
6 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (4464 kbps 14.9606466293 ms) (16.8490409851 ms 1.96981430054 ms 1.8907441101 ms) (391 390 8 pkt/s 2 %)
7 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (225 kbps 13.6257207 ms) (17.1720987327 ms 2.050742291 ms) (1711769104 ms) (868 856 4 pkt/s 0 %)
8 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (10562 kbps 29. packet_src, packet_dst, packet_loss, loss_rate (2.93207168579 ms) (916 912 4 pkt/s 0 %)
9 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (365 kbps 15.0351524353 ms) (17.2441095707 ms 2.1429061895 ms 2.27499008179 ms) (1011 1386 5 pkt/s 0 %)
10 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (16779 kbps 19.2122459412 ms) (20.768880441 ms 1.92022323608 ms 1.19304656987 ms) (1452 1447 5 pkt/s 0 %)
11 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (263 kbps 14.4288539886 ms) (17.303943634 ms 2.94804573059 ms 2.80213356018 ms) (999 999 0 pkt/s 0 %)
12 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (12132 kbps 24.4024991989 ms) (26.4630317681 ms 2.47502326965 ms 1.6400418701 ms) (1049 1046 3 pkt/s 0 %)
13 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (420 kbps 14.7120952660 ms) (16.5491104121 ms 1.8038749694 ms 1.87015533447 ms) (1594 1594 0 pkt/s 0 %)
14 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (19291 kbps 25.8771181107 ms) (27.1010398865 ms 1.50489807129 ms 0.942945480347 ms) (1664 1664 0 pkt/s 0 %)
15 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (285 kbps 16.2805318833 ms) (18.5420513153 ms 2.22992897034 ms 2.2931098938 ms) (1080 1080 0 pkt/s 0 %)
16 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (13139 kbps 24.9804258347 ms) (26.5829563141 ms 1.91042435303 ms 1.29103606585 ms) (1133 1133 0 pkt/s 0 %)
17 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (207 kbps 14.8888048172 ms) (17.7440643311 ms 3.05104225567 ms 2.66098976135 ms) (782 787 -5 pkt/s -1 %)
18 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (9506 kbps 16.1545276642 ms) (18.4030532837 ms 2.34794616699 ms 2.14910507202 ms) (815 820 -5 pkt/s -1 %)
19 s2-->s1 (00:00:00:00:00:02 - 00:00:00:00:00:03) - (289 kbps 12.6495361328 ms) (16.82305336 ms 4.45199012756 ms 3.89504432678 ms) (1097 1097 0 pkt/s 0 %)
20 s1-->s2 (00:00:00:00:00:01 - 00:00:00:00:00:02) - (13430 kbps 14.8576498032 ms) (18.1350708008 ms 3.37290763855 ms 3.18193435669 ms) (1155 1158 -3 pkt/s -1 %)

```

Figure 2.26: Loss 0 %, Linear Topology, tcp, server h1, client h2


```

1 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps 0.0 ms) (0 ms 0 ms 0 ms) (0 0 0 pkt/s 0 %)
2 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (0 kbps 0.0 ms) (0 ms 0 ms 0 ms) (0 0 0 pkt/s 0 %)
3 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (1 kbps -1.18279457092 ms) (3.96418571472 ms 5.72299957275 ms 4.57096099854 ms) (5 5 0 pkt/s 0 %)
4 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (31 kbps -0.592121541199 ms) (3.08585166931 ms 3.84283065796 ms 3.51309776306 ms) (14 10 4 pkt/s 28 %)
5 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps 0.0970363616943 ms) (5.9449672699 ms 5.126953125 ms 0.56890869141 ms) (0 0 0 pkt/s 0 %)
6 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (919 kbps -1.71637535095 ms) (2.8600692749 ms 5.93996047974 ms 3.21292877197 ms) (188 164 16 pkt/s 8 %)
7 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.217795372009 ms) (3.31616401672 ms 3.2958984375 ms 3.77202833997 ms) (0 0 0 pkt/s 0 %)
8 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (944 kbps -0.92339515686 ms) (1.66201591492 ms 3.26991081238 ms 1.9091133118 ms) (180 169 11 pkt/s 6 %)
9 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -1.12414360046 ms) (6.21294975281 ms 8.1741809845 ms 6.50000572265 ms) (0 0 0 pkt/s 0 %)
10 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (968 kbps -0.155925750732 ms) (6.03103637695 ms 6.2530049741 ms 6.12892018127 ms) (179 173 6 pkt/s 3 %)
11 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.803232192993 ms) (2.54988670349 ms 3.66115570068 ms 3.04588289229 ms) (0 0 0 pkt/s 0 %)
12 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (913 kbps -0.258564949036 ms) (2.1440982180 ms 2.45213508060 ms 2.35319137573 ms) (178 163 15 pkt/s 8 %)
13 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.1794099880774 ms) (3.55195999146 ms 3.37386131287 ms 4.08877863159 ms) (0 0 0 pkt/s 0 %)
14 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (996 kbps -1.03668351746 ms) (1.76087814331 ms 3.54504585266 ms 1.91807746887 ms) (184 178 6 pkt/s 3 %)
15 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.378489494324 ms) (2.77400016785 ms 3.08704376221 ms) (0 0 0 pkt/s 0 %)
16 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (953 kbps -0.744223594666 ms) (1.44982337952 ms 2.72703170776 ms 1.6610622406 ms) (181 170 11 pkt/s 6 %)
17 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.69415693054 ms) (2.93684005737 ms 3.65400314331 ms 3.6079835754 ms) (0 0 0 pkt/s 0 %)
18 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (953 kbps -0.85723400116 ms) (1.44982337952 ms 2.81000137329 ms 1.80411338806 ms) (181 170 11 pkt/s 6 %)
19 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -1.36649668612 ms) (7.33304023743 ms 9.45711135864 ms) (7.94196128845 ms) (0 0 0 pkt/s 0 %)
20 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (1041 kbps -0.827074050908 ms) (6.29997253418 ms 7.32010563416 ms 6.93392753601 ms) (192 185 7 pkt/s 3 %)
21 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.47492980597 ms) (2.12907791138 ms 2.76899337769 ms 2.43902206421 ms) (0 0 0 pkt/s 0 %)
22 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (980 kbps -0.210523605347 ms) (1.84798240662 ms 2.12597846985 ms 1.99103355408 ms) (181 175 6 pkt/s 3 %)
23 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.634074211121 ms) (4.84800338745 ms 5.79399463501 ms 5.17106056213 ms) (1 1 0 pkt/s 0 %)
24 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (831 kbps -0.668525695891 ms) (3.99804115295 ms 4.8451423645 ms 4.48799133301 ms) (160 150 10 pkt/s 6 %)
25 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (2 kbps 0.296950340271 ms) (4.40096855164 ms 3.25894355774 ms 4.94909286499 ms) (5 5 0 pkt/s 0 %)
26 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (2 kbps 0.862956047958 ms) (4.01306152344 ms 4.50420379639 ms 1.79600715637 ms) (7 6 1 pkt/s 14 %)
27 $5-->s1 (00:00:00:00:00:05 - 00:00:00:00:00:01) - (0 kbps -0.16725063324 ms) (3.56817245483 ms 3.41606140137 ms 4.05478477478 ms) (0 0 0 pkt/s 0 %)
28 $1-->s5 (00:00:00:00:00:01 - 00:00:00:00:00:05) - (0 kbps -0.15154514313 ms) (1.67107582092 ms 3.55410575867 ms 1.89113616943 ms) (0 0 0 pkt/s 0 %)

```

Figure 2.31: Loss 5 %, Complete Graph, udp, server h1, client h2

Figure 2.26 and Figure 2.28 compare test results of different loss rate configurations (0%, 5%) in TCP connection. Figure 2.29 to Figure 2.31 compare test results of different loss rate configurations and different topology configurations in UDP connection.

From the test results, we have noticed that some abnormal test results:

1. Negative loss rate: For example, in Figure 2.26, the configured loss rate is 0%, however the tested loss rate is below 0 for some test results.
2. Inaccurate loss rate: For example, in Figure 2.29, the configured loss rate is 5%, however, the tested loss rate is different from 5% for some test results.

We assume the reason for those abnormal results are caused by the reply statistics from the flow request messages. Once the **Monitoring** instance sends two flow request messages, two intermediate variables (**packet_src**, **packet_dst**) are returned by the reply message. **packet_src** is the number of packets sent by the server in the flow, **packet_dst** is the number of packets received by the client in the flow. However, when the **packet_dst** is returned by the reply message, there might exist some packets which haven't reached the destination switch.

2.5. Comparison

In this section, we will compare the advantages and disadvantages of an Active approach.

2.5.1. Advantages of Active approach

- It is less complex compared to passive approach, which gives real-time network statistics.
- This approach helps in providing network insights for a specified region.
- The data can be retrieved easily by polling the first and last switch.

2.5.2. Disadvantages of Active approach

- This approach does not measure real-time traffic delay, packet loss or bandwidth.
- It uses periodic polling is used, which gives less accurate results compared to adaptive polling.
- The probe packets generates additional traffic which can take up additional bandwidth and storage.

3

P4 Monitoring Implementation & Testing

3.1. P4 Monitoring

P4 (“Programming Protocol-independent Packet Processors”) is a language which specifies how packets are processed by the data plane of any switch, router or network interface card. In P4 also, there is a strict separation of data plane and control plane. These specifications can be installed on a switch and modified later. The operator can define how the packets are to be processed. In this way, P4 makes data plane flexible. The P4 monitoring implementation collects various throughput, delay and packet loss statistics.

This implementation of delay, bandwidth and loss monitoring is developed based on P4(v1.6). We will be making use of the provided virtual machine with the initial packet in packet out implementation of P4. In the report, we will mention, explain and screenshot most of the code we have added ourselves. For **delay monitoring**, we make uses of the ability to program the data plane to add information containing headers into the packets which in the end is forwarded to the controller and is extracted and processed by a monitoring thread. For **bandwidth monitoring** we will make use of the triple color meters and assign color rates in the central controller. For the **loss monitoring** we will use the integrated registers to contain the data which is sent to a monitoring thread in the controller. These three functionalities will be tested on a linear topology which consist of two hosts "h1", "h2" and three different switches "s1", "s2" and "s3", which are portrayed in Figure 3.1. Please notice that the controller is also portrayed and is part of the topology.

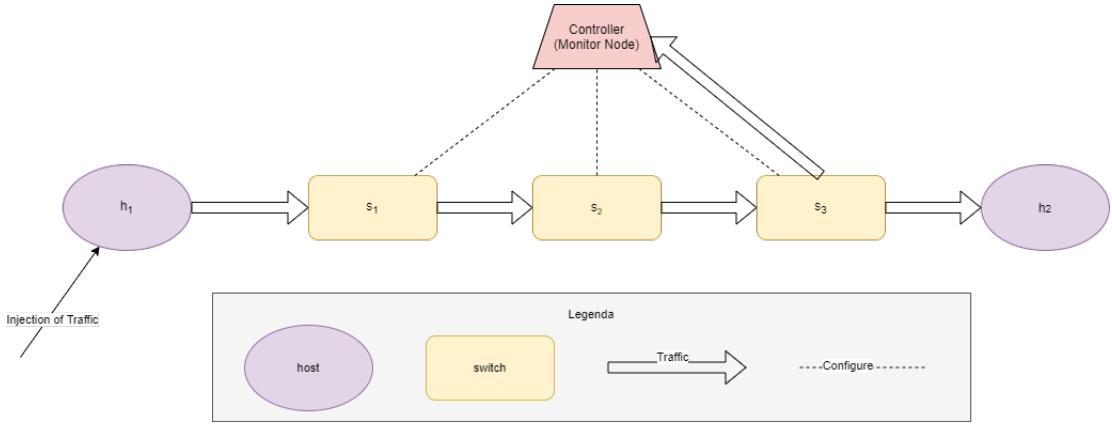


Figure 3.1: P4 testing topology 1

3.2. Diagnostic Node & Cloning

Before we begin with monitoring statistics from the network, a node is chosen to be the monitor node where monitoring information arrives and can be displayed. We have chosen the controller for various reasons: The controller is the "brain" of the network, therefore it is logical that statistics regarding the network should arrive at the controller. Another reason is, we can easily take advantage of the already existing implementation of packet in | packet out to forward packages containing information towards controller.

One of the most important thing to keep in mind is that, we need to implement this in such a way that the network traffic is unaffected. We choose to clone all packets containing network information from the final switch s3 before reaching destination h2 and send them towards controller. Imagine that network traffic from h1 to h2 would operate as normal and s3 is a "spy" node where all packets are forwarded towards the controller. Another approach could have been letting all the switches in the network clone and forward all packets towards the controller, but this approach is very impractical when it comes to a larger network where there are a lot of switches. With this approach, the controller would receive in a network of n nodes, per one single packet from h1 to h2, another extra n amount of cloned packets.

In order to perform cloning, we have to start a cloning session. This was done by initializing a cloning session in the controller.py¹ with the help for simple_switch_CLI. By default, SimpleSwitchGrpc do not support runtime_CLI, which needs to be enabled by compiling grpc switches with Thrift support [4].

After compiling the switch with Thrift support, we ran "netstat -tupnl" and got the following result, indicating success and all the switches are listening on their thrift ports.

¹controller.py is the file that contains the controller codes

```

File Edit View Search Terminal Help
hpdn@hpdn-VirtualBox:~$ netstat -tupnl
(Not all processes could be identified, non-owned process info
will not be shown, you would have to be root to see it all.)
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address          Foreign Address        State
tcp      0      0 127.0.0.53:53           0.0.0.0:*            LISTEN
tcp      0      0 127.0.0.1:631          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:9090          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:50050          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:9091          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:50051          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:9092          0.0.0.0:*            LISTEN
tcp      0      0 0.0.0.0:50052          0.0.0.0:*            LISTEN
udp      0      0 0.0.0.0:5353          0.0.0.0:*
udp      0      0 127.0.0.53:53          0.0.0.0:*
udp      0      0 0.0.0.0:68           0.0.0.0:*
udp      0      0 0.0.0.0:631          0.0.0.0:*
udp      0      0 0.0.0.0:33546          0.0.0.0:*
hpdn@hpdn-VirtualBox:~$ S

```

Figure 3.2: All thrift ports are listening

Now we set up the cloning. First, we used "Popen" from the subprocess module in the controller code to connect to switch 3 (which has thrift port of 9092) and created a mirroring session with the mirror id 2 (This was randomly chosen) and the egress port 255, which is the port towards the controller.

```

# Set up clone session for the final switch
print "\nSet up clone session from final switch towards controller"
s3_CLI = subprocess.Popen(["python behavioral-model/targets/simple_switch/simple_switch_CLI.in \
                           --thrift-port 9092"], shell = True, stdout = subprocess.PIPE, stdin = subprocess.PIPE)
s3_CLI.communicate(input = "mirroring_add 2 255")[0]

```

Figure 3.3: Setting up the a clone session in controller.py

Next, in the p4 code switch_final.p4² of the final switch, we used clone3 function to enable Ingress to Egress cloning [1], which create a new copy of the current packet from the end of the Ingress pipeline.

If the cloning is successful, we should receive ping replies from host 2 as well as receiving the new packets in the controller. See figure 3.4.

```

mininet> h1 ping h2
PING 10.10.10.2 (10.10.10.2) 56(84) bytes of data.
64 bytes from 10.10.10.2: icmp_seq=1 ttl=64 time=7.86 ms
64 bytes from 10.10.10.2: icmp_seq=2 ttl=64 time=4.48 ms
64 bytes from 10.10.10.2: icmp_seq=3 ttl=64 time=4.73 ms
64 bytes from 10.10.10.2: icmp_seq=4 ttl=64 time=3.67 ms
64 bytes from 10.10.10.2: icmp_seq=5 ttl=64 time=3.89 ms
64 bytes from 10.10.10.2: icmp_seq=6 ttl=64 time=7.51 ms
^C
--- 10.10.10.2 ping statistics ---
7 packets transmitted, 6 received, 14% packet loss, time 6036ms
rtt min/avg/max/mdev = 3.676/5.361/7.866/1.685 ms
mininet>

```

```

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-ma
File Edit View Search Terminal Help
-----PACKET IN received-----
Amount of packets received: 5
Packet size: 108
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)

-----PACKET IN received-----
Amount of packets received: 6
Packet size: 108
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)

```

Figure 3.4: Cloning successful. Both Host 2 and Controller has received the packets

²switch_final.p4 is the file containing code for the final switch in the path which is switch 3

3.3. Delay Monitoring

3.3.1. Implementation

The P4 delay monitoring is achieved by modifying the program of every switches to add an extra header called "latency_t" into each packet as instructed in the P4 lecture. The goal is to make sure that every hop in the network, the statistics can be added to an extra header. In this case, we would like to measure the queue delay of each switch, i.e. how long does a particular switch takes to process an average packet. This is essentials in the future in order to find out where in the network is the bottle neck and route traffic accordingly.

Within every P4 switch code in the egress we add a new action called "add_latency_information()" which is called every single time in the "apply" section of the egress. In this method we set the latency header to be valid as well as adding queuing delay of that particular switch by using information from the standard_metadata. We assign "queuing_delay" attribute from the header to be the "deq_timedelta" from the metadata, which is the time, in microseconds, that the current packet spent in the queue [5]. This code for the switches can be seen at switch.p4³. Do note that "deq_timedelta" will be 0 if assigned in the ingress, since the queuing take places between the ingress and egress. Therefore the correct way to get it is assign it in the egress.

The final switch on the path toward destination has the following unique function: besides adding delay information header like the other switches, it also clone the packet and forward them to both the original destination as well as the controller, as discussed in 3.2. A copy is made from the original packet which is forwarded to port 255 towards the controller. In the egress, we have implemented an if else case where if it is not a cloned packet, it is meant for the original recipients and all the extra latency headers should be removed. We do so by setting the headers stacks to invalid. The "latency_t" headers from the previous hops are removed so that this copy can be accepted by the destination as the "normal" packet. The destination will respond to the packet as if it was a normal packet without. This code for the final switch can be seen at "switch_final.p4"

³switch.p4 is the file containing code for the all switch in the network, except the final node which is switch 3

```

control MyEgress(inout headers hdr,
                 inout metadata meta,
                 inout standard_metadata_t standard_metadata) {
    action add_latency_information() {
        hdr.lastlatency.setValid();
        hdr.lastlatency.queuing_delay = standard_metadata.deq_timedelta;
    }
    apply {
        //Add switch latency information header, but only one way
        if(meta.stats_enabled){
            add_latency_information();
        }
        if(standard_metadata.instance_type == PKT_INSTANCE_TYPE_INGRESS_CLONE){
            //If this packet is a clone towards controller, we want to keep the headers for diagnostic purposes
            //So we do nothing
        } else{
            //If this packet is not a clone we want to remove the latency headers, to forward to original recipients as correct packet
            //Remove latency headers, but only one way
            if(meta.stats_enabled){
                //HARDCODED Element: there are 2 switches so there should be 2 latency headers
                hdr.latency[0].setInvalid();
                hdr.latency[1].setInvalid();
                hdr.lastlatency.setInvalid();
            }
        }
    }
}

```

Figure 3.5: Removing the latency headers in the egress if going towards original recipients

Notice that we hard-coded the amount of headers to be removed in the else case. This is not a good practice and should be avoided in the future by adding an extra variable in the header to indicate the final latency header when parsing the header stacks. An example implementing header stack and its parsing correctly so that there is no need for hard coding can be seen from a routing example [6].

Also notice that the "add_latency_information()" is only added one way, i.e. when the flow has stats enabled.

On the other hand, the clone is sent to the controller through the packet-in/packet-out mechanism. This copy still contains the delay headers which is extracted by a thread constantly running in the controller. For each packet according to the topology used, there will be 3 headers which belong to switch 1, 2 and 3 according to the order this packet went through in the network. These three headers are extracted and simply printed.

3.3.2. Testing

In order to test the delay monitoring, the first step is running a modified version of "UPDsend.py" provided by the course. This script can be seen at "LatencyTest_UPDsend.py" which will attempt to non-stop sending UDP packets with destination as Host 2. This script is run at Host 1 which can be accessed with "xterm h1" when running mininet.

```

#!/usr/bin/python

from scapy.all import *
import argparse
import random

def main():
    """
    """

    def packet_with_seq_n(port):
        packet = Ether(src="00:00:00:00:00:01", dst="00:00:00:00:00:02")/IP(src="10.10.10.1", dst="10.10.10.2")/UDP(sport=80, dport=port)/"UDPLatencyTest"
        sendp(packet, iface="eth0")

    if __name__ == "__main__":
        main()
        parser = argparse.ArgumentParser(description='Simple script that sends TCP packets to an interface using scapy package')
        parser.add_argument('-d', '--dst-port', help='Destination port', type=int, action='store', default=90)
        args = parser.parse_args()

        while True: #Send forever
            packet_with_seq_n(int(args.dst_port))

```

Figure 3.6: Code from the file LatencyTest_UDPsend.py

After running this "forever UDP sending" script, we take a look at the normal queue delay at the 3 switches which is parsed from the extra headers and printed on the CMD of the controller.py every time a packet arrives at the controller. The delay (microsecond!) is usually in the range of 2 to 3 digits at most.

```

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main
File Edit View Terminal Help
-----PACKET IN received-----
Amount of packets received: 769
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 67
Queue delay switch s2 (microsecond): 13
Queue delay switch s3 (microsecond): 62

-----PACKET IN received-----
Amount of packets received: 770
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 54
Queue delay switch s2 (microsecond): 17
Queue delay switch s3 (microsecond): 35

```

Figure 3.7: Queue delays in the range of 2 to 3 digits from three switches

Now we open up a simple_switch_CLI, connect to each of the switches 1, 2 and 3 with thrift ports 9090, 9091 and 9092 respectively and lower the queue rate with "set_queue_rate rate" command. See the following figures 3.8, 3.9 and 3.10 for the results.

```
hpdn@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main/behavioral-model
/targets/simple_switch$ python simple_switch_CLI.in --thrift-port 9090
Obtaining JSON from switch...
Done
Control utility for runtime P4 table manipulation
RuntimeCmd: set_queue_rate 20
RuntimeCmd: [REDACTED]
hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main
```

File Edit View Search Terminal Help

-----PACKET IN received-----
Amount of packets received: 13433
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3104796
Queue delay switch s2 (microsecond): 50
Queue delay switch s3 (microsecond): 27

-----PACKET IN received-----
Amount of packets received: 13434
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3100059
Queue delay switch s2 (microsecond): 105
Queue delay switch s3 (microsecond): 36

Figure 3.8: Delay increased after set queue rate of switch 1 to 20

```
hpdn@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main/behavioral-model
/targets/simple_switch$ python simple_switch_CLI.in --thrift-port 9091
Obtaining JSON from switch...
Done
Control utility for runtime P4 table manipulation
RuntimeCmd: set_queue_rate 20
RuntimeCmd: [REDACTED]
hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main
```

File Edit View Search Terminal Help

-----PACKET IN received-----
Amount of packets received: 15030
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3164973
Queue delay switch s2 (microsecond): 469653
Queue delay switch s3 (microsecond): 107

-----PACKET IN received-----
Amount of packets received: 15031
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3161705
Queue delay switch s2 (microsecond): 472571
Queue delay switch s3 (microsecond): 40

Figure 3.9: Delay switch 2 increased after set queue rate of switch 2 to 20

```

hpdn@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main/behavioral-model
/targets/simple_switch$ python simple_switch_CLI.in --thrift-port 9092
Obtaining JSON from switch...
Done
Control utility for runtime P4 table manipulation
RuntimeCmd: set_queue_rate 20
RuntimeCmd:
hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main
File Edit View Search Terminal Help
-----PACKET IN received-----
Amount of packets received: 17298
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3063368
Queue delay switch s2 (microsecond): 330417
Queue delay switch s3 (microsecond): 382776

-----PACKET IN received-----
Amount of packets received: 17299
Packet size: 66
Destination Address: (10, 10, 10, 1)
Source Address: (10, 10, 10, 2)
Queue delay switch s1 (microsecond): 3069581
Queue delay switch s2 (microsecond): 331038
Queue delay switch s3 (microsecond): 383700

```

Figure 3.10: Delay switch 3 increased after set queue rate of switch 3 to 20

The delays increase is caused due to the rate of processing packets in the queues of the switches are lowered. This is reflected in the tests.

We noticed that the delays started to increase dramatically after setting the queue rate near the point of 25 packets per second rate. Above 30 packets per seconds there are no changes to the delays experienced.

3.4. Bandwidth Monitoring

3.4.1. Implementation

The P4 bandwidth monitoring is achieved by implementing a meter which is one of the extern object that is supported by P4. The point of a meter is to measure the rate of flow of packets through a switch and classify them into one of the assigned three categories automatically. The high flows are classified as red, medium flows as yellow and low flows as green. They are classified based on the strength of the flow relative to each other in a meter which is an array of multiple states of multiple flows.

We start with creating a meter in switch.p4 file which is called "bandwidth_meter" with size of 10 and choose to measure packets with "Meter-Type.packets".

```
***** INGRESS PROCESSING *****/
control MyIngress(inout header$ hdr,
    inout metadata meta,
    inout standard_metadata_t standard_metadata) {
    //meter definition
    meter(10, MeterType.packets) bandwidth_meter; //measuring based on packets
```

Figure 3.11: Meter initialization in switch.p4

On the headers we have also defined three different meter level representing green, yellow and red rate of packet flow. They are represented with 2 bit field "meter_level_t" as integer 0, 1 and 2 respectively, see 3.12 . We also added a variable "meter_level" into the meta data to hold the status of the flow.

```
//params for meter
typedef bit<2> meter_level_t;
const meter_level_t METER_COLOR_GREEN = 0;
const meter_level_t METER_COLOR_YELLOW = 1;
const meter_level_t METER_COLOR_RED = 2;
```

Figure 3.12: Meter colors initialization in switch.p4

In the ingress, we have implemented two different tables in order to make the meter work properly. The first table is called "flow_table" which has the function of classifying new packages to the right "index", which is used to identify one of the states of the meter. This is done by matching the packet's destination MAC address, i.e. the packets belong to the same flow/index/state of a meter if they have the same destination MAC address. We have implemented this table in such a way that in the controller, you can customize which switch to perform the bandwidth monitoring with the meter and which destination MAC address to keep track off.

The second table in the Ingress is called "flow_filter_table" which serves the purpose of filtering the packets by their meter level. For example, in the controller.py, we can install rules such that all packets belongs to a flow that has "become" red, i.e. become to large in the amount of incoming packets, to be dropped. For instance, if a packet comes in the "meta.meter_level" is 2, we can install the drop action.

```
//Table to keep tracks of different flow according to rules added in the controller.py
table flow_table{
    key = {
        hdr.ethernet.dstAddr: exact;
    }
    actions = {
        modify_flow_level;
        NoAction;
    }
}
size = 1024;
default_action = NoAction();

//Table to filter and drop flows according to rules added in the controller.py
table flow_filter_table{
    key = {
        meta.meter_level : exact;
    }
    actions = {
        drop;
        NoAction;
    }
}
size = 16;
default_action = NoAction();
```

Figure 3.13: Implementation of flow_table and flow_filter_table in switch.py

For the implementation of the meter and the tables, refer to the file "switch.p4". For the implementation of writing the rules into these tables, refer to the file "controller.py". In the controller, we only write methods/ rules for switch 1, which is to measure all traffic with destination mac address to be host 2 and only drop level 1 packets, see 3.14. Since we have added tables to the switch code, we had to modify the "firmeware.p4info.txt" as well to include the new tables.

```
# Writing rules
# Forwarding rules
writeIpv4Rules(p4info_helper, sw_id=s1, dst_ip_addr="10.10.10.1", port = 1)
# withStats = True to only include latency header one waySS
writeIpv4Rules(p4info_helper, sw_id=s1, dst_ip_addr="10.10.10.2", port = 2, withStats = True)

# Rule for bandwidth monitoring on switch 1
# Add which destination mac address to add a meter state to monitor
writeFlowTableRules(p4info_helper, sw_id=s1, dst_mac_addr="00:00:00:00:00:02")
# Add which at meter rules in order to drop yellow (level 1) packets
#writeFlowFilterTableRules(p4info_helper, sw_id=s1, drop_level= 1)

readTableRules(p4info_helper, s1) #read all table rules
```

Figure 3.14: Meter assignment and bandwidth drop rules in controller.py

3.4.2. Testing

To test the bandwidth monitoring function, we start with running mininet as usual and then run controller.py to add rules to the switches. This time, start two XTERM windows of host 1 and host 2 with "xterm h1" and "xterm h2" respectively. On host 2 we start a UDP listen with "iperf -s -u -i 1" and on host 1 we sent UDP data with "iperf -c 10.10.10.2 -u -b 1M -l 1000". We are sending UDP datagrams from host 1 to host 2 with bandwidth of 1 megabit

per second and datagram size of 1000 bytes. See 3.15 below, you can see that the bandwidth towards host 2 was roughly 1 megabit per second as desired.

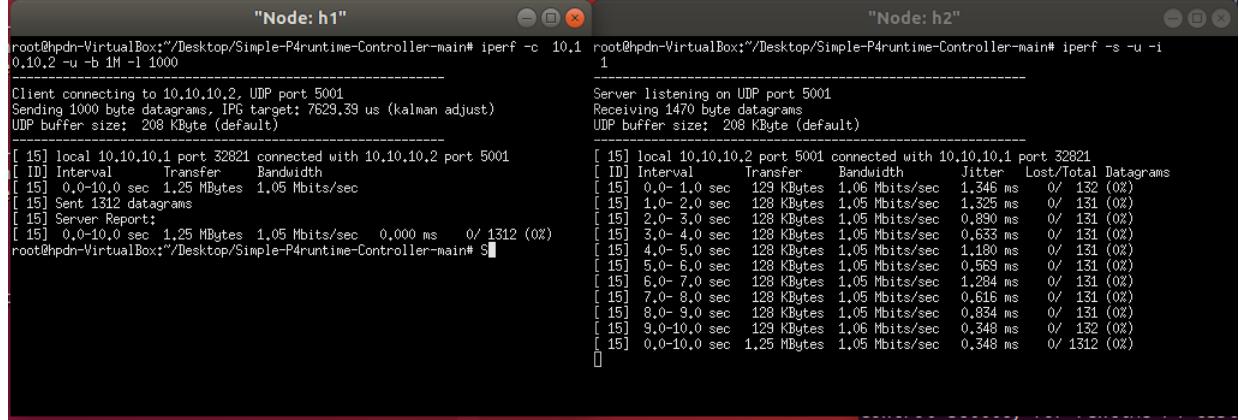


Figure 3.15: Meter assignment and bandwidth drop rules in controller.py

Now we will open up the simple_switch_CLI once again, connect with switch 1 and run the following command "meter_set_rates" bandwidth_meter 2 0.0001:1 0.0005:1". This means that we set the rate for the yellow color of the meter to be 100 packets per seconds.

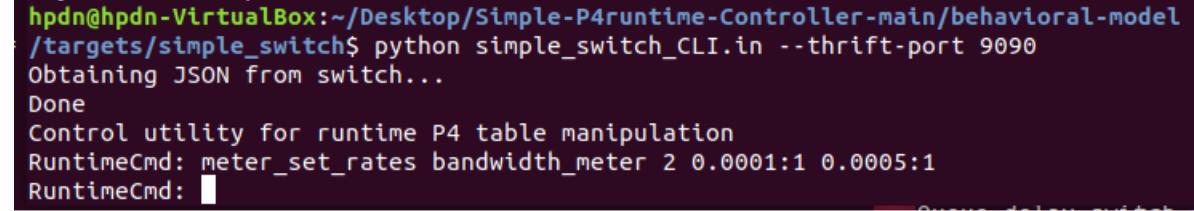


Figure 3.16: Setting rates for the meter in switch 1

Do remember from figure 3.14 that we have added rules to drop all packets that are yellow (drop_level = 1). So the maximum bandwidth achieved after dropping packets should be equal $100 \text{ packets/ second} * 1000 \text{ bytes/ packet} = 100000 \text{ bytes/ second}$. This is equal to 800 Kbits/second. The following figure shows the result of running the same iperf test several time again.

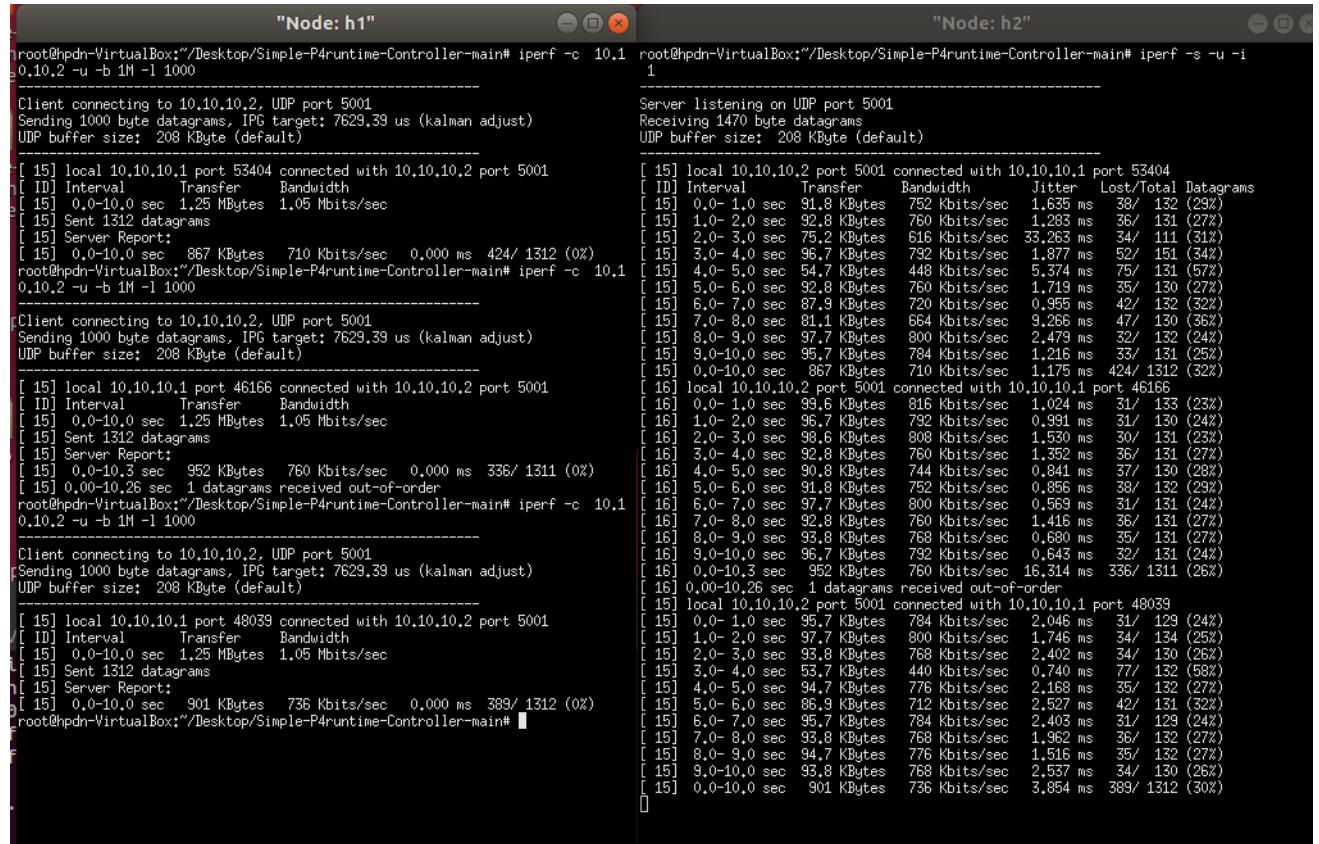


Figure 3.17: Result of rerunning the iperf test several times again

As you can see, the maximum bandwidth after setting the meter rates with dropping rules are always roughly 800 Kbits/second. This means the meter is working properly.

As a bonus, we disabled the rule to drop level 1 packets in the controller, rerun the controller and rerun the iperf test. Indeed from the result below you can see that the packets are not dropped anymore and the maximum bandwidth is back to roughly 1 Mbit/second.

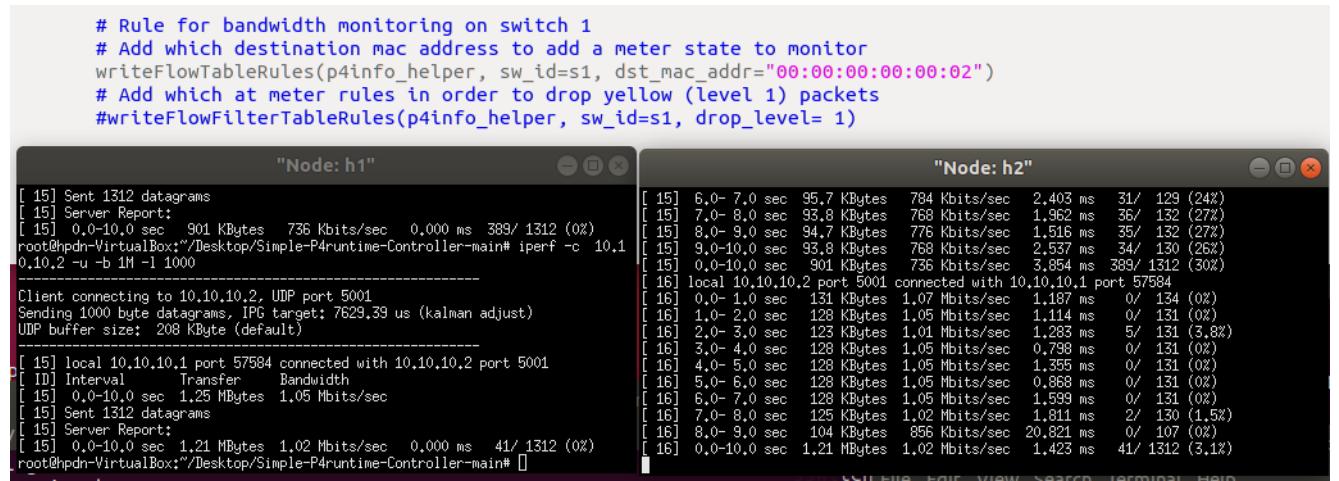


Figure 3.18: Bandwidth return back to normal after disable the meter dropping rule in controller.py

3.5. Loss Monitoring

3.5.1. Implementation

Loss monitoring was achieved with the help of using registers. In both switch.p4 and switch_final.p4, we have added a register called packet_counter. This register is used to log the amount of packets received based on the IPv4 identification header. The idea is that if a packet fragment consisting of 100 packets (which have the same IPv4 identification header) passes through a switch, the register with the index of that identification number will be incremented 100 times, if there are no packet loss at that switch. We have an action called update_counter which read the register at the index which is equal to IPv4 identification number of a packet and increment it. See figure 3.19.

```

control MyEgress(inout headers hdr,
                  inout metadata meta,
                  inout standard_metadata_t standard_metadata) {

    //register definition for loss monitoring
    register<bit<32>>(65536) packet_counter;

    action update_counter(bit<16> id){

        //First read the amount counted packets with this id
        bit<32> count;
        packet_counter.read(count, (bit<32>)id);

        //Increment by 1
        count = count + 1;
        //record the ID into the register
        packet_counter.write((bit<32>)id, count);
    }

    //action for delay calculation
    action add_latency_information() {
        hdr.latency.setValid();
        hdr.latency.queuing_delay = standard_metadata.deq_timedelta;
    }

    apply {
        update_counter(hdr.ipv4.identification); //update for loss monitoring
        if(meta.state.enabled)
    }
}

```

Figure 3.19: Implementation of register and method to update it in the switch.p4

In retrospect, this can also be done with a counter, but we have decided to keep register in case of further expansion of loss monitoring in the future.

3.5.2. Testing

To test loss monitoring, we would run from Host 1 a python file called LossTest_UDPsend.py. This file would send in total 1000 UDP packets with a payload consisting of 2 things: the total number of packets to be received and the number of the current packet in sequence. The former variable is called n and the later is called seq_number. We have randomly chosen a IPv4 identification number to be 112 for all packets. See figure 3.20.

```
#!/usr/bin/python

from scapy.all import *
import argparse
import random
import struct

def main():
    """
    """

    def packet_with_seq_n(port, n, seq_number):
        packet = Ether(src="00:00:00:00:00:01", dst="00:00:00:00:00:02")/IP(src="10.10.10.1", dst="10.10.10.2", id = 112)/UDP(sport=7777, dport=port)/Raw(load= struct.pack("ii", n, seq_number))
        sendp(packet, iface="eth0")

    if __name__ == "__main__":
        main()
        parser = argparse.ArgumentParser(description='Simple script that sends TCP packets to an interface using scapy package')
        parser.add_argument('-dst-port', help='Destination port', type=int, action="store", default=90)
        args = parser.parse_args()

        n =1000 #Total number of packets to be sent
        for seq_number in range(1, n+1):
            packet_with_seq_n(int(args.dst_port), n, seq_number)
```

Figure 3.20: Implementation of LossTest_UDPsend.py

In the controller.py we have written code to extract n and seq_number (mentioned previously) of each packets and make statistics out of it, e.g. the total number of packets not received by the controller and are lost. We also used CLI to access each registers of each switch in order to fetch the number of packets counts from the same identification number of 112. See figure 3.21.

```
#####Code for TESTING LOSS MONITORING
total = 0;
current = 0;
print "\nReady for loss test, please run LossTest_UDPsend.py from Host 1"
while True:
    packetin = s3.PacketIn()           #Packet in!
    if packetin is not None:
        received += 1
        print "\n-----PACKET IN received-----"
        payload = packetin.packet.payload

    (total, current) = concatAndUnpack(payload, 'ii', 52, 59)
    print "Loss monitoring in progress: %f %%" % (current/float(total) *100)
    print "Total number of packets to be received: %d " % total
    print "Total number of packets actually received: %d " % received

    if((current/float(total) *100) == 100): #end of test
        break;
s| receivedAtSwitch1 = subprocess.Popen(["python behavioral-model/targets/simple_switch/simple_switch_CLI.in \
                                         -thrift-port 9090"], shell = True, stdout = subprocess.PIPE, stdin = subprocess.PIPE).communicate(input = "register_read packet_counter 112")[0]
    receivedAtSwitch2 = subprocess.Popen(["python behavioral-model/targets/simple_switch/simple_switch_CLI.in \
                                         -thrift-port 9091"], shell = True, stdout = subprocess.PIPE, stdin = subprocess.PIPE).communicate(input = "register_read packet_counter 112")[0]
    receivedAtSwitch3 = subprocess.Popen(["python behavioral-model/targets/simple_switch/simple_switch_CLI.in \
                                         -thrift-port 9092"], shell = True, stdout = subprocess.PIPE, stdin = subprocess.PIPE).communicate(input = "register_read packet_counter 112")[0]

    receivedAtSwitch1 = int(filter(str.isdigit, receivedAtSwitch1[118:150]))
    receivedAtSwitch2 = int(filter(str.isdigit, receivedAtSwitch2[118:150]))
    receivedAtSwitch3 = int(filter(str.isdigit, receivedAtSwitch3[118:150]))

    print "\n\n***** SUMMARY LOSS MONITORING STATISTICS *****"
    print "Total number of packets to be received: %d " % total
    print "Total number of packets actually received by controller: %d " % received
    print "Total number of packets lost on the way (not received by controller): %d " % (total-received)

    print "Number of packets lost at SWITCH 1: %d " % (total-receivedAtSwitch1)
    print "Number of packets lost at SWITCH 2: %d " % (total-receivedAtSwitch2)
    print "Number of packets lost at SWITCH 3: %d " % (total-receivedAtSwitch3)
    print "Loss percentage at SWITCH 1: %f %%" % (((total-receivedAtSwitch1)/float(total))*100)
    print "Loss percentage at SWITCH 2: %f %%" % (((total-receivedAtSwitch2)/float(total))*100)
    print "Loss percentage at SWITCH 3: %f %%" % (((total-receivedAtSwitch3)/float(total))*100)
    print "***** END OF SUMMARY *****"
```

Figure 3.21: Implementation of loss monitoring testing in the controller.py

Now its time to test loss monitoring. First we run a control test with normal settings with sending 1000 UDP packets and capture the result. See figure 3.22

The figure shows three terminal windows on a Linux desktop. The top-left window is titled "Node: h1" and displays a loop of "Sent 1 packets." messages. The top-right window shows the command being run: "hpdn@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/targets/simple_switch\$ python simple_switch_CLI.py --thrift-port 9092". The bottom window shows the controller's log output:

```

h1 h2
*** Starting controller
*** Starting 3 switches
s1 Starting P4 gRPC switch s1
simple_switch_grpc -i 1@s1-eth1 -i 2@s1-eth2 --thrift-port 9090 --device-id 1 switch.json --log-flush --log-level trace --log-file s1.log -- --grpc-server-addr 0.0.0.0:50050 --cpu-port 255
switch has been started successfully
s2 Starting P4 gRPC switch s2
simple_switch_grpc -i 1@s2-eth1 -i 2@s2-eth2 --thrift-port 9091 --device-id 2 switch.json --log-flush --log-level trace --log-file s2.log -- --grpc-server-addr 0.0.0.0:50051 --cpu-port 255
switch has been started successfully
s3 Starting P4 gRPC switch s3
simple_switch_grpc -i 1@s3-eth1 -i 2@s3-eth2 --thrift-port 9092 --device-id 3 switch_final.json --log-flush --log-level trace --log-file s3.log -- --grpc-server-addr 0.0.0.0:50052 --cpu-port 255
switch has been started successfully
Starting mininet
*** Starting CLI:
mininet> xterm h1
mininet>

```

The rightmost terminal window shows the loss monitoring summary:

```

-----PACKET IN received-----
Loss monitoring in progress: 99.900000 %
Total number of packets to be received: 1000
Total number of packets actually received: 999

-----PACKET IN received-----
Loss monitoring in progress: 100.000000 %
Total number of packets to be received: 1000
Total number of packets actually received: 1000

***** SUMMARY LOSS MONITORING STATISTICS *****
Total number of packets to be received: 1000
Total number of packets actually received by controller: 1000
Total number of packets lost on the way (not received by controller): 0
Number of packets lost at SWITCH 1: 0
Number of packets lost at SWITCH 2: 0
Number of packets lost at SWITCH 3: 0
Loss percentage at SWITCH 1: 0.000000 %
Loss percentage at SWITCH 2: 0.000000 %
Loss percentage at SWITCH 3: 0.000000 %
***** END OF SUMMARY *****


```

Figure 3.22: Control test

As you can see, as already expected and can be seen in a summary of loss monitoring statistics, there would be no packet loss at any switch under normal condition.

Now we introduce an artificial packet loss by simply increasing the queue processing rate of switch 2 with "set_queue_rate 10" and see what happens. Note that we only set the queue rate to 10 for only a couple of seconds before we set it back to maximum queue rate. See figure 3.23.

The screenshot shows three terminal windows on a Linux desktop. The top-left window is titled "Node: h1" and displays a loop of "Sent 1 packets." messages. The top-right window is titled "hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/t..." and shows a command-line interface (CLI) session. The CLI outputs JSON configuration for a switch, followed by runtime commands like "set_queue_rate 10" and "set_queue_rate 1000". The bottom-left window is titled "File Edit View Search Terminal Help" and shows the startup logs for three switches (s1, s2, s3) and a mininet node (h1). The bottom-right window is titled "hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main" and displays the results of a loss monitoring test. It shows statistics for each switch: s1 has 99.9% loss (1 packet lost), s2 has 100% loss (1 packet lost), and s3 has 22.8% loss (228 packets lost). The total number of packets sent was 1000, and 771 were received.

```

"Node: h1"
Sent 1 packets.
.
.
.
root@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main# 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ python simple_switch_CLI.in --thrift-port 9091
Obtaining JSON from switch...
Done
Control utility for runtime P4 table manipulation
RuntimeCmd: set_queue_rate 10
RuntimeCmd: set_queue_rate 1000
RuntimeCmd: 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ 

File Edit View Search Terminal Help
h1 h2
*** Starting controller
*** Starting 3 switches
s1 Starting P4 gRPC switch s1
simple_switch_grpc -i 1@s1-eth1 -i 2@s1-eth2 --thrift-port 9090 --device-id 1 switch.json --log-flush --log-level trace --log-file s1.log -- --grpc-server-addr 0.0.0.0:50050 --cpu-port 255
switch has been started successfully
s2 Starting P4 gRPC switch s2
simple_switch_grpc -i 1@s2-eth1 -i 2@s2-eth2 --thrift-port 9091 --device-id 2 switch.json --log-flush --log-level trace --log-file s2.log -- --grpc-server-addr 0.0.0.0:50051 --cpu-port 255
switch has been started successfully
s3 Starting P4 gRPC switch s3
simple_switch_grpc -i 1@s3-eth1 -i 2@s3-eth2 --thrift-port 9092 --device-id 3 switch_final.json --log-flush --log-level trace --log-file s3.log -- --grpc-server-addr 0.0.0.0:50052 --cpu-port 255
switch has been started successfully
Starting mininet!
*** Starting CLI:
mininet> xterm h1
mininet> 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ python simple_switch_CLI.in --thrift-port 9091
File Edit View Search Terminal Help
-----PACKET IN received-----
Loss monitoring in progress: 99.900000 %
Total number of packets to be received: 1000
Total number of packets actually received: 771
-----PACKET IN received-----
Loss monitoring in progress: 100.000000 %
Total number of packets to be received: 1000
Total number of packets actually received: 772
***** SUMMARY LOSS MONITORING STATISTICS *****
Total number of packets to be received: 1000
Total number of packets actually received by controller: 772
Total number of packets lost on the way (not received by controller): 228
Number of packets lost at SWITCH 1: 0
Number of packets lost at SWITCH 2: 228
Number of packets lost at SWITCH 3: 228
Loss percentage at SWITCH 1: 0.000000 %
Loss percentage at SWITCH 2: 22.800000 %
Loss percentage at SWITCH 3: 22.800000 %
***** END OF SUMMARY *****

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ 

```

Figure 3.23: Loss test with induced loss at switch 2

In the summary you can see that there are 228 packets NOT dropped at switch 2 and therefore they did not arrive at switch 3 as well. This number matches the amount of packets that is NOT received by the controller. The loss monitoring test has correctly displayed the amount and percentage of loss packets in each switch. See figure 3.24 to see the result of a test where a lowered queue rate is temporary induced at switch 3.

The figure shows three terminal windows on a Linux desktop. The top-left window, titled "Node: h1", displays a loop of packet transmission: "Sent 1 packets." followed by a colon and a dot. The top-right window, titled "hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/tar...)", shows a Python script running. It includes code for the `sswitch_CLI.main()` function, which handles keyboard input and performs configuration via Thrift port 9092. The bottom-left window, titled "hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/tar...)", shows the CLI interface. It starts a controller on host 1 (h1) and three switches (s1, s2, s3). The bottom-right window, also titled "hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/tar...)", displays the results of a loss monitoring test. It shows that 99.9% of 1000 packets were received at switch 3, with 125 packets lost.

```

"Node: h1"
Sent 1 packets.
Sent 1 packets.
.
Sent 1 packets.
root@hpdn-VirtualBox:/Desktop/Simple-P4runtime-Controller-main# 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main/behavioral-model/tar... 
File Edit View Search Terminal Help
sswitch_CLI.main()
  File "/home/hpdn/Desktop/Simple-P4runtime-Controller-main/behavioral-model/tar...
gets/simple_switch/sswitch_CLI.py", line 144, in main
  SimpleSwitchAPI(args.pre, standard_client, mc_client, sswitch_client).cmdloo...
p()
  File "/usr/lib/python2.7/cmd.py", line 138, in cmdloop
    line = raw_input(self.prompt)
KeyboardInterrupt
hpdn@hpdn-VirtualBox:~/Desktop/Simple-P4runtime-Controller-main/behavioral-model...
/tar.../simple_switch$ python simple_switch_CLI.in --thrift-port 9092
Obtaining JSON from switch...
Done
Control utility for runtime P4 table manipulation
RuntimeCmd: set_queue_rate 10
RuntimeCmd: set_queue_rate 1000
RuntimeCmd: 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main
File Edit View Search Terminal Help
h1 h2
*** Starting controller

*** Starting 3 switches
s1 Starting P4 gRPC switch s1
simple_switch_grpc -i 1@s1-eth1 -i 2@s1-eth2 --thrift-port 9090 --d...
evice-id 1 switch.json --log-flush --log-level trace --log-file s1.log ...
--grpc-server-addr 0.0.0.0:50050 --cpu-port 255
switch has been started successfully
s2 Starting P4 gRPC switch s2
simple_switch_grpc -i 1@s2-eth1 -i 2@s2-eth2 --thrift-port 9091 --d...
evice-id 2 switch.json --log-flush --log-level trace --log-file s2.log ...
--grpc-server-addr 0.0.0.0:50051 --cpu-port 255
switch has been started successfully
s3 Starting P4 gRPC switch s3
simple_switch_grpc -i 1@s3-eth1 -i 2@s3-eth2 --thrift-port 9092 --d...
evice-id 3 switch_final.json --log-flush --log-level trace --log-fil...
le s3.log -- --grpc-server-addr 0.0.0.0:50052 --cpu-port 255
switch has been started successfully

Starting mininet!
*** Starting CLI:
mininet> xterm h1
mininet> 

hpdn@hpdn-VirtualBox: ~/Desktop/Simple-P4runtime-Controller-main$ 

```

Figure 3.24: Loss test with induced loss at switch 3

As expected of correct loss monitoring, switch 3 is showed here to have dropped/not received 125 packets out of 1000 packets.

4

Results & Comparison

4.1. Openflow

4.1.1. Bandwidth monitoring

S. No.	Configured Bandwidth (Mbps)	Maximum Throughput (Mbps)	Topology	Figure
1.	20	19.606	Linear with two switches.	Fig 2.5
2.	20	18.5	Linear with five switches.	Fig 2.6
3.	20	19.2	Linear with five switches along with two clients and hosts	Fig 2.7
4.	40	38.1	Linear with two switches.	Fig 2.9
5.	40	37.5	Linear with five switches	Fig 2.10
6.	40	38.6	Linear with five switches along with two clients and hosts	Fig 2.11
7.	20	19.606	Ring with five hosts and five switches	Fig 2.8
8.	40	38.1	Ring with five hosts and five switches	Fig 2.12

Table 4.1: Bandwidth Results

- From Fig 2.6-2.7, it is evident that there are two flows, one from switch s2 to switch s1 and the second one is from s1 to s2. The size of the first flow is smaller compared to the second flow because the second flow is carrying the payload while the first flow consists of ACK messages as explained in CS4055 High Performance Data Networking Reader. The gradual increase in throughput demonstrates the TCP slow start algorithm which is a mechanism for congestion control. The initial sender's window size is smaller compared to receiver's window size. However, it doubles the rate every RTT until congestion is detected.
- From the Table 4.1, it can be clearly inferred that the maximum throughput which was measured by the monitoring app was close to the configured bandwidth but these two values were never equal. This observation stands true for all kinds of topologies and bandwidth configurations that were tested. The reason behind the marginal differences in the throughput value are latency, jitter and packet loss.
- The throughput also decreased as the number of switches were increased from two to five as observed in 2 and 5. After adding three more switches

between the source and destination, the latency increased which reduced the maximum throughput.

- Further, if we compare ring topology and linear topology (from 1,7 and 6,8) the maximum throughput remains the same because in our case ring topology acts as point-to-point linear topology.

4.1.2. Delay Monitoring

- The measured delay was found to be in milliseconds for every topology which was tested upon.
- The delay monitoring app was tested on different topologies such as linear and complete graph with varying nodes. The delay results were same for linear as well as complete mesh topology. The reasons for the unexpected delays were overhead due to Linux kernel and stacking of intermediate variables used for delay calculation as discussed in chapter 2.

4.1.3. Loss Monitoring

- The packet loss was greater than the assigned value of 5 percent with TCP traffic as observed in Fig 2.28 and equal to 0 percent for assigned value of 0 percent(Fig 2.26).
- For UDP traffic, the packet loss was higher than the assigned value and also higher when compared to the TCP traffic for the same assigned loss rate.
- The reason for higher packet loss is that UDP does not contain any congestion control mechanism (sliding window) which is present in TCP. As a result, the packets start dropping if there is congestion and UDP keeps sending the packets at the same rate.
- Another observation was that the bandwidth was also reduced due to the increase in packet loss percentage.

4.2. P4

4.2.1. Bandwidth monitoring

- The throughput achieved was equal to the assigned bandwidth when the rule to drop level 1 packet was disabled which was 1 Mbit per second. However, after configuring the rule to drop level 1 packet the achieved throughput got reduced to 800 Kbits per second.
- By using meters, the network bandwidth management becomes easier. If a flow becomes red, which means the rate of the packets arriving at the flow exceeds the specified rate and drop packet action is configured then the switch will start dropping packets which acts as a means of congestion control. In this manner, a better user experience is ensured.
- The meter also safeguards the network from packets of high data rates, that might hog the entire bandwidth of the network by limiting the rate of each flow.

- From Fig 3.15 and Fig 3.17, we can notice that since we have sent UDP traffic, the observed throughput was constant after uniform time intervals unlike TCP which used slow start algorithm for congestion control.
- Currently, we have limited the number of packets for each flow but this functionality can also be extended to limit the data rate which exceeds the specified limit.

4.2.2. Delay Monitoring

- The queuing delay was found to be in milliseconds, given in Fig 3.7. The queuing delay is the time a packet waits in the queue before getting processed.
- The delay monitoring app measured the correct queuing delay of all three switches. This delay information can be used to detect which node in the network is taking longer time to process the packets. Then, bottlenecks in the network can be avoided by finding out the root cause, for example replacing an outdated switch that has very slow processing time. These delays also affect the network throughput.
- Delay Monitoring can be used in the future to find out if a particular path is slow and switch to other paths that have faster processing switches.

4.2.3. Loss Monitoring

- From Fig 3.22, we can observe that the number of lost packets were 0, which means that there were no network issues such as bandwidth unavailability, jitters and increased latency.
- From Fig 3.23, decrease in queue processing rate resulted in small packet loss, as it introduced a larger queuing delay in the network.
- If there is a packet loss at one switch(s1) then the next few switches will also suffer packet loss because there will be no forwarding from the switch S1.

4.3. OpenFlow and P4 Comparison

4.3.1. Openflow monitoring approach

- For measuring network statistics, OpenFlow follows active approach which means injecting additional packets for collecting network statistics.
- The drawback of this approach is that it does not give real-time network statistics. In addition, these probe packets generate additional traffic which takes up extra bandwidth and storage.
- The polling interval should be small for getting updated statistics from time to time. Even if the polling frequency is in seconds, microsecond congestion cannot be avoided.

4.3.2. P4 monitoring approach

- The P4 monitoring follows a different approach. The monitoring app works by creating additional headers which are encoded in the data packet itself.
- The drawback of this approach is the increased packet size due to extra headers.
- The increase in packet size can be a problem for networks which consist of large number of nodes as these will consume more bandwidth.
- Also cloning was done to replicate packets which were sent towards the controller which acted as the monitoring node and aggregated network statistics. Together these cloned packets and additional headers can take up significant bandwidth.

From above, it follows that both these discussed approaches have various advantages and disadvantages. The P4 approach is more efficient for smaller networks consisting of fewer number of nodes.

4.3.3. Comparison based on Overhead

- For OpenFlow bandwidth monitoring and loss monitoring, first and last switch were regularly polled by sending stats request packets. The polling frequency was chosen as 2 seconds. Similarly, additional probe packet was created to measure the end-to-end delay. These packets are responsible for the extra overhead.
- For delay monitoring, a custom header was created for aggregating the latency statistics per hop. For bandwidth monitoring, configuring the meter with the required rate was enough for bandwidth management along with addition of header for meter in the beginning of the code. For loss monitoring, register was used to keep track of the data fragments and later these statistics were collected from the monitoring node.
- Consider the latency monitoring in which the data packet became 32 bit long after the last hop. This result is shown in fig This increase in packet size will not be suitable for networks which consists of larger number of nodes since with every hop the packet size will considerably increase.

```

    }

header latency_t{
    bit<32> queuing_delay;
}

```

Figure 4.1: latency header

From above observations, it can be inferred that OpenFlow overhead is greater as compared to the overheads in P4. Instead of creating additional probe packets, telemetry information is embedded inside the header which is less compared to the OpenFlow overhead.

4.3.4. Comparison based on Network Statistics

The P4 and OpenFlow are also compared based on the following parameters:

- **Bandwidth monitoring:** In P4 the throughput was equal to the configured bandwidth whereas in OpenFlow the measured throughput was always less than the set limit.
- **Loss monitoring:** The packet loss measured was similar for both P4 as well as OpenFlow.
- **Delay monitoring:** In P4 delay monitoring, we have measured only queueing delay and not the end-to-end delay between the switches because queueing delay information is useful in locating the switch which has gone out of order. This can be used in future to find the congestion bottleneck. In Openflow, end-to-end delay was measured which takes into account propagation delay, transmission delay and queueing delay.

5

Conclusion

The goal of this project was to do a comparative study of OpenFlow monitoring application with P4 monitoring application. This comparison was based on bandwidth, delay and packet loss. For carrying out the comparison, we first developed applications both for P4 as well as OpenFlow for monitoring network statistics. Then we analysed the results from each one of these apps and compared those results with iperf results for validation. The approach followed for collecting data in OpenFlow and P4 were different from each other. After implementing these monitoring functionalities, both the apps were thoroughly tested with varying traffic, topologies and link bandwidths.

Based on these test results, it can be concluded that P4 is more efficient compared to the OpenFlow as the overhead is less, it provides flexible data plane, reduces the load on controller to a great extent and achieved maximum throughput. However, for monitoring in large-sized networks, especially with latency monitoring, the overhead might increase due to additional headers.

Bibliography

- [1] Randeng works. URL <https://carolinafernandez.github.io/development/2019/08/06/Recurrent-processing-in-P4>.
- [2] Gary Lee. *Cloud networking: Understanding cloud-based data center networks*. Morgan Kaufmann, 2014.
- [3] D. Marschke, J. Doyle, and P. Moyer. *Software Defined Networking (SDN): Anatomy of OpenFlow Volume I*. Number v. 1. Lulu Publishing Services, 2015. ISBN 9781483427249. URL <https://books.google.co.in/books?id=ecM5CgAAQBAJ>.
- [4] p4lang. p4lang/behavioral-model,. URL https://github.com/p4lang/behavioral-model/tree/main/targets/simple_switch_grpc#enabling-the-thrift-server.
- [5] p4lang. p4lang/behavioral-model,. URL https://github.com/p4lang/behavioral-model/blob/main/docs/simple_switch.md.
- [6] p4lang. p4lang/tutorials,. URL https://github.com/p4lang/tutorials/blob/master/exercises/source_routing/solution/source_routing.p4.
- [7] Niels LM Van Adrichem, Christian Doerr, and Fernando A Kuipers. Opennetmon: Network monitoring in openflow software-defined networks. In *2014 IEEE Network Operations and Management Symposium (NOMS)*, pages 1–8. IEEE, 2014.